

ZHAW SCHOOL OF ENGINEERING
INSTITUTE OF MECHATRONIC SYSTEMS

Betaflight – Offline Controller Tuning

Data-Driven Tuning of Rate and Altitude Controllers for FPV Drones

Bachelor Thesis

Yuri Bianchi, Dario Jurietti and Janick Dort

Supervisor: Michael Peter (pmic)

Co-Supervisor: Prof. Dr. Ruprecht Altenburger (altb)

March 31, 2026

Declaration of Originality

We hereby declare that this thesis is our own work and has been composed solely by the undersigned. We have not used any sources or aids other than those indicated. All passages taken verbatim or in substance from published or unpublished works of others have been clearly marked as such. This thesis has not been submitted, in whole or in part, for any other degree or examination at this or any other institution.

Winterthur, March 31, 2026

Yuri Bianchi

Dario Jurietti

Janick Dort

Zusammenfassung

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Abstract

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Contents

Declaration of Originality	ii
Zusammenfassung	iii
Abstract	iv
List of Figures	vi
List of Tables	vii
Preface	viii
1 Introduction	1
1.1 Initial situation	2
1.2 Task and Objective	2
2 Theoretical Foundations	3
2.1 Chirp Signal	3
2.1.1 Mathematical Definition	3
2.1.2 Common Frequency Profiles	4
2.1.3 Lag Filter for Chirp Input Shaping	6
2.2 Signal Processing	7
2.2.1 Baseband Filtering via Frequency Rotation	9
2.2.2 Spectral Estimation using Welch's Method	11
2.3 Control Loop Calculations	14
2.3.1 Calculation of the Control Loop Components	14
2.3.2 Closed-Loop Performance Metrics	15
3 Methodology	17
3.1 Angle Mode	17
3.1.1 Difference between Angle and Acro Mode	17
3.1.2 Control Loops in Angle Mode	17
3.1.3 Derivation of the Target Angle	18
3.1.4 Chirp Signal for System Identification	19
Declaration of Generative AI Use	20

List of Figures

2.1	Linear chirp signal used for the simulation.	4
2.2	Exponential chirp signal used for the simulation.	5
2.3	Exponential chirp signal after conditioning with the lag filter used for chirp input shaping.	6
2.4	Simulation model used for the signal processing analysis.	7
2.5	Bode plot of the true and estimated transfer function. Deviations between the curves occur in frequency regions dominated by noise and disturbances. . . .	9
2.6	Bode plot of the true transfer function and the estimated transfer function based on the filtered signal. The filtering improves the signal-to-noise ratio and leads to a more accurate estimation.	11
2.7	Bode plot of the true and estimated transfer function. A good agreement between both curves indicates that the spectral estimation captures the system dynamics well.	13
2.8	Magnitude-squared coherence between input and output signals. Values close to one indicate a strong linear relationship and therefore a reliable transfer function estimate, whereas low values indicate reduced reliability due to noise and disturbances.	14
3.1	Cascaded control structure of the Betaflight Angle mode controller.	18
3.2	Current control loop structure in Betaflight when the chirp signal is injected inside the closed loop.	19

List of Tables

Preface

This thesis was carried out during the spring semester 2026 at the Institute of Mechatronic Systems (IMS) at the ZHAW School of Engineering. It continues the work started in the preceding project thesis conducted in the fall semester 2025.

We would like to thank our supervisor, Michael Peter, for...

1 Introduction

First-Person-View (FPV) racing drones require precise and responsive control to achieve stable flight behaviour, suppress disturbances such as propeller wash, and maintain agility during rapid manoeuvres. The flight performance strongly depends on the tuning of the flight controller's feedback loops. In particular, the parameters of the angular rate controller determine how quickly and accurately the drone reacts to pilot inputs and external disturbances.

In practice, controller tuning is still often performed manually through iterative trial-and-error procedures. This requires extensive flight testing and considerable piloting experience, making the process time-consuming and difficult to reproduce. Although modern flight controller firmware such as Betaflight provides extensive configuration options and detailed flight data logging, systematic tuning remains a challenge.

One commonly used tool for offline tuning is the PID-Toolbox. In this approach, flight logs are recorded during normal flights and later analysed to derive suitable controller parameters. However, the system excitation depends entirely on the pilot's inputs and the natural flight behaviour of the drone. As a result, not all relevant frequency ranges are sufficiently excited, which can limit the quality of the measurements.

In the approach proposed in this thesis, a chirp excitation signal is injected directly into the control loop. This allows the system to be excited over a wide range of frequencies in a controlled and systematic way, resulting in higher-quality measurements for system identification and controller tuning.

A previous project introduced a proof-of-concept approach for offline tuning of angular rate controllers based on recorded flight data. In the preceding project thesis, this concept was restructured and modularised and now serves as the foundation for the present work.

The objective of this bachelor thesis is to further refine and extend this approach into an offline tuning framework for multiple control loops used in Betaflight-based flight controllers. In addition to the existing angular rate controller tuner, the framework will be extended with code modules for tuning the angle controller, the altitude hold controller, and the position controller. This extends the tuning framework beyond the inner rate control loop and enables the analysis and tuning of higher-level control layers.

The framework will be implemented primarily in MATLAB and Python and will provide tools for analysing flight logs, estimating system dynamics, and generating controller parameters without requiring repeated flight tests. The software will be developed as an open-source library with accompanying documentation for the FPV community.

By providing a structured, data-driven approach to controller tuning, this project aims to simplify the tuning process and improve flight performance.

1.1 Initial situation

1.2 Task and Objective

2 Theoretical Foundations

In this chapter, the theoretical background of the approach used for tuning the different control loops of the flight controller is introduced. Since the fundamental working principle is identical for all loops, the general concepts are explained here before addressing the individual loops in the methodology section.

To support the theoretical concepts, a simplified simulation model is introduced. The simulation is gradually developed throughout this chapter, allowing the reader to follow the modelling process step by step and to better understand the behaviour of the control system. This approach illustrates the influence of the controller parameters and provides an intuitive understanding of the tuning process.

The detailed application of these principles to the different control loops implemented in Betaflight is discussed in the following chapter.

2.1 Chirp Signal

2.1.1 Mathematical Definition

The instantaneous value of the chirp signal is defined as

$$x(t) = \sin(\arg(t)),$$

where the phase argument

$$\arg(t) = 2\pi \int_0^t f(\tau) d\tau$$

represents the instantaneous phase of the signal. In practical implementations such as **Betaflight**, the phase is wrapped once it reaches 2π . This keeps the numerical values bounded and produces the sawtooth-shaped **sinarg** signal that is often observed in practice.

The behaviour of the chirp signal is determined by several parameters that define the frequency sweep:

- f_0 – starting frequency of the chirp signal (in the simulation 1 Hz)
- f_1 – final frequency of the chirp signal (in the simulation 600 Hz)
- T – total duration of the frequency sweep (in the simulation 20 s)
- f_s – sampling frequency used in the discrete-time simulation (in the simulation 2000 Hz)

The parameters used in the simulation are chosen to be close to the typical operating range of the real system.

2.1.2 Common Frequency Profiles

Linear chirp:

For a linear chirp, the instantaneous frequency increases linearly over time. The slope of the frequency sweep is defined as

$$k = \frac{f_1 - f_0}{T}.$$

The instantaneous frequency is therefore given by

$$f(t) = f_0 + kt,$$

and the corresponding phase trajectory results from integrating the frequency over time

$$\arg(t) = 2\pi \left(f_0 t + \frac{1}{2} k t^2 \right).$$

Figure 2.1 illustrates the resulting chirp signal, the instantaneous frequency $f(t)$, and the corresponding phase trajectory $\arg(t)$ for the simulation parameters introduced in the previous section.




Figure 2.1: Linear chirp signal used for the simulation.

The apparent gaps visible in the chirp signal and in the wrapped phase are caused by the finite sampling time of the simulation and the limited resolution of the plotted data, rather than by discontinuities in the signal itself.

Exponential chirp (as used in Betaflight):

For an exponential chirp, the instantaneous frequency increases exponentially from the starting frequency f_0 to the final frequency f_1 over the time interval T . The instantaneous frequency is defined as

$$f(t) = f_0 \left(\frac{f_1}{f_0} \right)^{t/T}.$$

The corresponding phase trajectory results from integrating the frequency over time

$$\arg(t) = \frac{2\pi T f_0}{\ln(f_1/f_0)} \left[\left(\frac{f_1}{f_0} \right)^{t/T} - 1 \right].$$

These expressions define the phase trajectory $\arg(t)$ used in the simulations and in the `apply_rotfiltfilt` method for time-dependent frequency shifting of chirp signals. The resulting exponential chirp signal, the instantaneous frequency, and the phase trajectory are illustrated in Figure 2.2.




Figure 2.2: Exponential chirp signal used for the simulation.

The apparent gaps in the chirp signal and the wrapped phase result from the discrete sampling time of the simulation and the limited resolution of the plot.

2.1.3 Lag Filter for Chirp Input Shaping

Before the chirp signal is injected into the control loop, it is conditioned by a filter to address the differentiating behavior of the controller effort at low frequencies. This differentiating characteristic causes rising gain at low frequencies, which would lead to excessive motor response at the start of the chirp sweep and potentially cause saturation, compromising the quality of system identification data. The filter reduces low-to-mid frequency components of the chirp signal before injection, such that the combined response produces approximately uniform excitation across the frequency range.

The filter is configured with a pole at 3 Hz and a zero at 30 Hz. This configuration corresponds to the transfer function

$$H(s) = \frac{1 + s/\omega_z}{1 + s/\omega_p} \quad (2.1)$$

where $\omega_z = 2\pi \cdot 30$ rad/s and $\omega_p = 2\pi \cdot 3$ rad/s represent the zero and pole frequencies respectively. The pole introduces the lag component by creating magnitude reduction at frequencies above its cutoff frequency. The zero limits this reduction by introducing magnitude amplification at frequencies above its cutoff frequency. However, since the pole frequency is lower than the zero frequency, the pole's contribution dominates throughout the primary control bandwidth, resulting in a net lag filter. It is important to note that in the Betaflight implementation, the pole frequency (ω_p) must not be set higher than the zero frequency (ω_z), as this would create a lead filter that amplifies high frequencies and could potentially damage the motors.

For the simulation shown in Figure 2.3, the same parameter values introduced in the previous section were used ($f_0 = 1$ Hz, $f_1 = 600$ Hz, $T = 20$ s, $f_s = 2000$ Hz). The excitation signal is generated using the exponential chirp formulation described earlier.




Figure 2.3: Exponential chirp signal after conditioning with the lag filter used for chirp input shaping.

2.2 Signal Processing

To illustrate the influence of the applied signal processing steps, a simulation is used as a reference. The simulated system approximates the dynamics of the drone attitude control system and consists of an integrator, a first-order lag element, and a dead-time component.

The resulting transfer function is

$$G(s) = G_1(s) G_2(s) G_t(s)$$

with

$$G_1(s) = \frac{\omega_1}{s}, \quad G_2(s) = \frac{1}{1 + s/\omega_2}, \quad G_t(s) = e^{-s/\omega_t}.$$

Thus, the overall transfer function becomes

$$G(s) = \frac{\omega_1}{s} \cdot \frac{1}{1 + s/\omega_2} \cdot e^{-s/\omega_t}.$$

For the simulation, the parameters are chosen as

$$\omega_1 = 16 \cdot 2\pi, \quad \omega_2 = 13 \cdot 2\pi, \quad \omega_t = 60 \cdot 2\pi.$$

The structure of the simulation model is illustrated in Figure 2.4. The chirp signal $u(t)$ is applied to the transfer function $G(s)$, resulting in the ideal system output $\tilde{y}(t)$. To emulate realistic measurement conditions, additive disturbance noise $d(t)$ is superimposed on the output, such that

$$y(t) = \tilde{y}(t) + d(t).$$

In the simulation, the disturbance is modelled as additive white Gaussian noise with a standard deviation of $\sigma_d = 2 \text{ deg/s}$.




Figure 2.4: Simulation model used for the signal processing analysis.

2.2.0.1 Estimate Frequency Response

For a meaningful estimation of the transfer function, the frequency response between the input and output signals is analyzed.

In the frequency domain, the relation between the input and output of a linear time-invariant (LTI) system can be expressed as

$$G(\omega) = \frac{Y(\omega)}{U(\omega)}.$$

Here, $U(\omega)$ and $Y(\omega)$ denote the Fourier transforms of the input signal $u(t)$ and the output signal $y(t)$, respectively. This ratio describes how each frequency component of the input is scaled and phase-shifted by the system.

In practice, however, the measured output signal is affected by disturbances. The measured output can therefore be written as

$$Y(\omega) = G(\omega)U(\omega) + D_y(\omega),$$

where $D_y(\omega)$ represents an additive disturbance at the output. If the transfer function were estimated directly by dividing $Y(\omega)$ by $U(\omega)$, the disturbance term would directly influence the result

$$\frac{Y(\omega)}{U(\omega)} = G(\omega) + \frac{D_y(\omega)}{U(\omega)}.$$

This expression shows that a direct estimation is highly sensitive to disturbances and measurement noise.

To mitigate this issue, it is advantageous to consider the relation between input and output in terms of spectral quantities. In a first step, spectral quantities can be formed directly from the Fourier transforms as

$$S_{uu}(\omega) = |U(\omega)|^2, \quad S_{yu}(\omega) = Y(\omega)U^*(\omega).$$

Based on these quantities, the transfer function can be expressed as

$$G(\omega) = \frac{S_{yu}(\omega)}{S_{uu}(\omega)}.$$

However, this formulation is equivalent to the direct division $Y(\omega)/U(\omega)$ and therefore inherits the same sensitivity to disturbances. In particular, the disturbance term is fully contained in the cross-spectral quantity and directly propagates into the estimate, which may lead to significant deviations from the true system behavior.

A more robust formulation is obtained by defining the spectral densities in a statistical sense using expectation values. This reduces the influence of disturbances, provided that the noise is uncorrelated with the input signal. Since the expectation operator cannot be evaluated directly from finite measurement data, it must be approximated by averaging over multiple

realizations or segments of the signals. In practice, this is achieved using spectral estimation techniques such as Welch's method, which is introduced in the following section.

As shown in Fig. 2.7, the estimated transfer function closely follows the true system response in frequency regions where the signal-to-noise ratio is sufficiently high. In contrast, deviations occur in frequency ranges dominated by noise and disturbances, where the estimation becomes less reliable.




Figure 2.5: Bode plot of the true and estimated transfer function. Deviations between the curves occur in frequency regions dominated by noise and disturbances.

2.2.1 Baseband Filtering via Frequency Rotation

Measured signals are often contaminated by measurement noise and unwanted high-frequency components. These disturbances can obscure the relevant information contained in the signal and negatively affect subsequent analysis steps. To suppress these disturbances while preserving the phase characteristics of the signal, a frequency-shifting filtering approach is employed.

The method is based on shifting the frequency content of the signal to baseband, applying a low-pass filter, and subsequently shifting the signal back to its original frequency range. This allows the filtering operation to focus on a narrow frequency region around the signal of interest.

The underlying principle can be understood using the Fourier shift theorem. Multiplication of a signal $f(t)$ with a complex exponential results in a shift of its spectrum

$$e^{i2\pi\xi_0 t} f(t) \iff \hat{f}(\xi - \xi_0),$$

which translates the spectrum by the frequency ξ_0 .

In the present case, the signal of interest is a chirp signal with a continuously varying instantaneous frequency. Therefore, the frequency shift must also be time-dependent. This is achieved by introducing a complex phasor

$$p(t) = e^{i\phi(t)},$$

where $\phi(t)$ denotes the instantaneous phase of the chirp signal.

Multiplying the measured signal $x(t)$ with this phasor or its complex conjugate results in a rotation of the signal spectrum toward baseband

$$y_R(t) = x(t)p(t), \quad y_Q(t) = x(t)p^*(t).$$

After this rotation step, the dominant signal component is centered around 0 Hz, enabling efficient filtering using a low-pass filter.

2.2.1.1 Baseband Filtering

Once the signal has been shifted to baseband, a low-pass filter is applied to remove high-frequency noise components. Since the useful signal is concentrated near zero frequency, the filter can effectively suppress unwanted spectral components outside this region.

To avoid phase distortion, forward-backward filtering is employed. This approach applies the filter in both forward and reverse directions, resulting in zero-phase filtering.

The filtering operation yields the filtered baseband signals

$$y_R(t) \rightarrow \tilde{y}_R(t), \quad y_Q(t) \rightarrow \tilde{y}_Q(t),$$

which contain the desired signal components with reduced noise.

2.2.1.2 Inverse Frequency Rotation

After filtering in the baseband, the signal is shifted back to its original frequency range. This is achieved by multiplying the filtered signals with the corresponding inverse phasors and combining the results to obtain a real-valued signal

$$x_f(t) = \text{Re} \left(\frac{1}{2} (\tilde{y}_R(t)p^*(t) + \tilde{y}_Q(t)p(t)) \right).$$

This operation reverses the initial frequency rotation and reconstructs a filtered version of the original signal. The resulting signal $x_f(t)$ retains the phase characteristics of the original signal while high-frequency noise components are effectively suppressed.

2.2.1.3 Impact on Frequency Response Estimation

The filtered signal is subsequently used for frequency response estimation. By reducing noise prior to the spectral analysis, the signal-to-noise ratio is significantly improved, leading to more reliable estimates of the transfer function.

The effect of the filtering approach is illustrated in Fig. 2.6, where the estimated transfer function based on the filtered signal is compared with the true system response. A closer agreement between both curves indicates that the baseband filtering effectively suppresses noise while preserving the relevant system dynamics.




Figure 2.6: Bode plot of the true transfer function and the estimated transfer function based on the filtered signal. The filtering improves the signal-to-noise ratio and leads to a more accurate estimation.

2.2.2 Spectral Estimation using Welch's Method

To estimate the spectral densities $S_{yu}(\omega)$ and $S_{uu}(\omega)$ from finite data, Welch's method is employed. The spectral densities cannot be computed directly from the expectation operator, since only finite time-domain measurements of the input and output signals are available. Instead, these quantities must be estimated from the measured data. In this work, Welch's method is used to obtain reliable estimates of the power spectral densities and cross-power spectral densities.

Let $u[n]$ and $y[n]$ denote the discrete-time input and output signals of length N . In Welch's method, the signals are divided into K overlapping segments of length L . Each segment is multiplied by a window function $w[n]$ in order to reduce spectral leakage.

The k -th windowed segment of the input signal can therefore be written as

$$u_k[n] = u[n + kD] w[n], \quad n = 0, \dots, L - 1$$

where D denotes the shift between successive segments and determines the amount of overlap between segments.

For each segment, the discrete Fourier transform (DFT) is computed

$$U_k(\omega) = \sum_{n=0}^{L-1} u_k[n]e^{-j\omega n}$$

and analogously for the output signal

$$Y_k(\omega) = \sum_{n=0}^{L-1} y_k[n]e^{-j\omega n}.$$

From these Fourier transforms, the periodogram estimate of the power spectral density of the input signal for each segment is obtained as

$$P_{uu,k}(\omega) = \frac{1}{L}|U_k(\omega)|^2.$$

Similarly, the cross-periodogram between output and input is defined as

$$P_{yu,k}(\omega) = \frac{1}{L}Y_k(\omega)U_k^*(\omega).$$

Welch's method reduces the variance of the spectral estimates by averaging the periodograms over all segments. The resulting estimates of the power spectral density and the cross-power spectral density are therefore given by

$$S_{uu}(\omega) = \frac{1}{K} \sum_{k=1}^K P_{uu,k}(\omega)$$

and

$$S_{yu}(\omega) = \frac{1}{K} \sum_{k=1}^K P_{yu,k}(\omega).$$

In addition, the power spectral density of the output signal is obtained as

$$S_{yy}(\omega) = \frac{1}{K} \sum_{k=1}^K P_{yy,k}(\omega),$$

with

$$P_{yy,k}(\omega) = \frac{1}{L} |Y_k(\omega)|^2.$$

These averaged spectral estimates are then used to compute the transfer function estimate

$$G(\omega) = \frac{S_{yu}(\omega)}{S_{uu}(\omega)}.$$

Furthermore, the coherence function can be defined as

$$\gamma^2(\omega) = \frac{|S_{yu}(\omega)|^2}{S_{uu}(\omega) S_{yy}(\omega)}.$$

The coherence takes values in the range $0 \leq \gamma^2(\omega) \leq 1$ and provides a frequency-dependent measure of the linear relationship between input and output signals. Values close to one indicate that the output is predominantly determined by the input, whereas low values indicate that noise, disturbances, or nonlinear effects dominate the measurement and reduce the reliability of the estimated transfer function.

2.2.2.1 Estimated Transfer Function

The estimated transfer function obtained from the averaged spectral quantities is compared with the true system response in the frequency domain, as shown in Fig. 2.7. A good agreement between both curves indicates that the spectral estimation provides a reliable approximation of the underlying system dynamics. Deviations occur primarily in frequency regions where the signal-to-noise ratio is low and the estimation becomes less accurate.




Figure 2.7: Bode plot of the true and estimated transfer function. A good agreement between both curves indicates that the spectral estimation captures the system dynamics well.

In order to further assess the reliability of the estimated transfer function, the coherence function is evaluated. Unlike the transfer function itself, the coherence provides a direct measure of how strongly the output signal is linearly related to the input signal at each frequency.

2.2.2.2 Coherence Analysis

The coherence corresponding to the estimated transfer function is shown in Fig. 2.8. High coherence values indicate frequency regions in which the output is strongly driven by the input signal and the transfer function estimate can be considered reliable. In contrast, low coherence values indicate that noise, disturbances, or nonlinear effects dominate the measurement, leading to increased uncertainty in the estimated transfer function.




Figure 2.8: Magnitude-squared coherence between input and output signals. Values close to one indicate a strong linear relationship and therefore a reliable transfer function estimate, whereas low values indicate reduced reliability due to noise and disturbances.

2.3 Control Loop Calculations

Although the flight controller consists of several nested control loops, their fundamental structure is identical. Each loop follows the same basic feedback control principle and differs mainly in the controlled variable and the chosen controller parameters. Therefore, instead of analysing every loop individually from the beginning, a simplified model is used to explain the underlying calculations and controller behaviour.

The simplified control structure used for the analysis consists of a controller C and a plant P , which describes the dynamic behaviour of the system being controlled.

Using this simplified model allows the general behaviour of the control loop and the calculation of the controller terms to be explained in a clear and structured way. The same principles are later applied to the individual control loops implemented in the flight controller.

2.3.1 Calculation of the Control Loop Components

After the signal processing steps, the control loop components can be calculated. As in the signal processing stage, the goal is to mitigate the influence of disturbances and noise on the

calculated control loop components. Therefore, only the input signal of the control loop is used for the estimation.

The advantage of this approach is that

- the input signal is not affected by the back-calculation of the control loop and is therefore not influenced by disturbances or noise within the loop.
- the input signal is not affected by filtering steps, since it is not measured but taken directly from the software. Consequently, it is not subject to measurement noise or disturbances.

This approach allows a more precise estimation of the transfer function. However, it also means that the transfer functions of the individual control loop components cannot be calculated directly. Instead, they must be obtained through mathematical manipulations.

For example, consider a simplified control loop consisting of a controller $C(s)$ and a plant $P(s)$.

To calculate the transfer function of the controller $C(s)$, the following relation can be used

$$C(s) = \frac{Y(s)}{E(s)} = \frac{Y(s)}{R(s) - Y(s)} = \frac{Y(s)}{R(s)} \cdot \frac{1}{1 - \frac{Y(s)}{R(s)}} = \frac{G(s)}{1 - G(s)}.$$

To calculate the transfer function of the plant $P(s)$, the following relation can be used

$$P(s) = \frac{Y(s)}{U(s)} = \frac{Y(s)}{E(s)} \cdot \frac{E(s)}{U(s)} = C(s) \cdot \frac{1}{C(s)} = 1.$$

2.3.2 Closed-Loop Performance Metrics

After determining the plant transfer function, the next step is to analyze the behaviour of the overall system when a controller is applied. This is achieved by evaluating the closed-loop transfer functions, which describe the system response, disturbance sensitivity, disturbance rejection capability, and the required controller effort. These measures provide valuable insights into the stability, robustness, and performance of the controlled system.

The closed-loop transfer function from the reference input to the output is given by

$$G_{yr}(\omega) = \frac{C(\omega)P(\omega)}{1 + C(\omega)P(\omega)}.$$

The sensitivity function, which describes how sensitive the system output is to disturbances and model uncertainties, is defined as

$$S(\omega) = \frac{1}{1 + C(\omega)P(\omega)}.$$

The compliance function, representing the system response to disturbances acting on the plant, is given by

$$SP(\omega) = P(\omega) S(\omega).$$

Finally, the controller effort describes the influence of the reference input on the controller output and is defined as

$$SC(\omega) = C(\omega) S(\omega).$$

2.3.2.1 Step Response Analysis

To be added.

3 Methodology

In this chapter, the practical application of the theoretical concepts introduced in the previous chapter is described. The focus is on how these principles are implemented in the context of tuning the control loops of a Betaflight-based flight controller. The methodology is structured around the different control loops, starting with the angle controller, followed by the altitude hold controller and the position controller. For each control loop, the specific implementation details and tuning procedures are described.

3.1 Angle Mode

3.1.1 Difference between Angle and Acro Mode

The Betaflight software provides three main flight modes: Angle, Horizon, and Acro.

In Angle mode, the drone's tilt angle is directly linked to the stick position. The maximum tilt is limited to a predefined value, preventing flips and making this mode particularly suitable for beginners.

In contrast, Acro mode provides full manual control. In this mode, stick inputs define an angular rate rather than a target angle, meaning the drone continues rotating as long as the stick remains deflected. This allows aggressive maneuvers such as flips and rolls but also requires greater pilot experience.

During the project thesis, an offline tuner for the angular rate controller was further developed to systematically optimize the rate control loop and improve Acro mode performance. The tuner was subsequently extended to also support Angle mode, enabling improved stability and responsiveness while maintaining intuitive handling for beginner pilots.

3.1.2 Control Loops in Angle Mode

In Angle mode, the angular rate controller is extended by an outer angle control loop, resulting in a cascaded control structure.

The stick input initially generates a rate setpoint. Unlike in Acro mode, this rate setpoint is then converted into a target angle. This target angle serves as the reference input for the inner angular rate controller.

The cascaded architecture separates attitude control from angular rate stabilization. While the outer loop ensures convergence toward the desired orientation, the inner loop provides fast dynamic stabilization by regulating the angular velocity. As a result, the system achieves improved disturbance rejection and more intuitive handling characteristics, since the pilot directly commands an angle rather than a rotational rate.

The Angle mode controller therefore consists of two cascaded loops:

- The outer angle loop, which computes the angle error
- The inner angular rate loop, which computes the rate error




Figure 3.1: Cascaded control structure of the Betaflight Angle mode controller.

The angle control loop consists of a proportional controller with an additional feedforward term. Within the offline tuning framework, however, only the proportional controller is optimized. The feedforward path is included in the control structure for completeness but was deliberately disabled during the tuning process.

3.1.3 Derivation of the Target Angle

The target angle is obtained by scaling the commanded angular rate with respect to the maximum achievable rate and the configured maximum tilt angle

$$\theta_{\text{target}} = \theta_{\text{limit}} \cdot \frac{\omega_{\text{set}}}{\omega_{\text{max}}} \quad (0.1)$$

where

- θ_{limit} is the maximum allowed tilt angle (e.g., 60°),
- ω_{set} is the commanded angular rate derived from the stick input,
- ω_{max} is the maximum achievable angular rate determined by the rate configuration.

This relation ensures a proportional mapping between stick deflection and target angle.

For example,

$$\omega_{\text{set}} = \omega_{\text{max}} \Rightarrow \theta_{\text{target}} = \theta_{\text{limit}} \quad (0.2)$$

$$\omega_{\text{set}} = 0 \Rightarrow \theta_{\text{target}} = 0 \quad (0.3)$$

Thus, the stick position scales the target angle continuously within the range

$$-\theta_{\text{limit}} \leq \theta_{\text{target}} \leq +\theta_{\text{limit}} \quad (0.4)$$

3.1.4 Chirp Signal for System Identification

As already mentioned in the gyro control tuning section, a chirp signal is injected into the control loop.

If the chirp signal were added in the same way as for the gyro control tuning, it would not be possible to directly calculate the transfer function of the complete system from the target angle to the measured angle. The chirp signal would be added to the current setpoint and therefore inside the closed control loop. Consequently, the system would have two inputs: the stick input and the injected chirp signal.

This situation is illustrated in Figure 3.2, which shows the control loop structure currently implemented in the Betaflight firmware.




Figure 3.2: Current control loop structure in Betaflight when the chirp signal is injected inside the closed loop.

To overcome this issue, a small modification in the firmware allows the chirp signal to be added directly to the setpoint before it is converted into the target angle. Due to the transformation from the setpoint to the target angle, the parameters of the chirp signal do not need to be adjusted.

The modified control structure used for the offline tuning framework is illustrated in Figure ?? . In this configuration, the chirp signal is injected before the transformation from the rate setpoint to the target angle, allowing the transfer function from the target angle to the measured angle to be estimated directly.

Some older drones are only capable of flying in Acro mode and therefore do not support Angle mode. In this case, the chirp signal can still be added to the setpoint. Since the angle control loop and the transformation from the setpoint to the target angle are disabled in Acro mode, the resulting control loop structure is identical to the one used for the gyro control tuning.

This approach allows the same chirp signal to be used for both modes. As a result, the implementation is simplified and the tuning process becomes more straightforward, since the same excitation signal can be applied without considering structural differences in the control loop.

Declaration of Generative AI Use

The following generative AI systems were used during the preparation of this thesis. Their use was limited to the functions described below. All generated content was critically reviewed and verified by the authors.